

CSE 415 – Intro to Parallel Computing
Spring 2021, Homework 2
Due 11:59 pm, Friday, Feb 12th

Important note: Please use a word processing software (e.g., MS Word, Mac Pages, Latex, etc.) to type your homework. Follow the submission instructions at the end to turn in an electronic copy of your work.

1) [25 pts] Performance Modeling

Assume you are evaluating a polynomial of the form, $x = y^2 + z^3 + yz$, implemented as follows:

```
float x[N], y[N], z[N];
for (i=0; i < N; ++i)
    x[i] = y[i]*y[i] + z[i]*z[i]*z[i] + y[i]*z[i];
```

Here i is an integer and x, y and z are single precision floating point arrays.

- a. [10 pts] What is the arithmetic intensity of this kernel?
- b. [15 pts] Develop a roofline model for a system with a memory bandwidth of 30 GB/sec, and a peak computational throughput of 95 GFLOP/sec. Mark the arithmetic intensity of the kernel on your plot and determine its expected performance. Is the performance memory-bound or compute-bound?

2) [75 pts] Cache optimization: Matrix Vector Block Multiplication

In this problem, you will optimize the multiplication of an $N \times M$ matrix with block of vectors (whose width for purposes of this assignment is taken to be constant at $N_{VEC}=16$). This routine takes as input a vector block of size M by N_{VEC} , and produces an output vector block of size N by N_{VEC} . A simple nested loop implementation is given below.

```
for (i = 0; i < N; i++) {
    for (j = 0; j < M; j++) {
        for (k = 0; k < NVEC; k++) {
            output[i][k] += input_matrix[i][j]*input_vector[j][k];
        }
    }
}
```

This simple implementation cannot make efficient use of the memory hierarchy. In particular, assuming row-major storage, accesses to the input vector block may result in a high number of cache misses.

- a. [25 pts] *Implementation:* Develop a cache optimized version of the matrix vector multiplication kernel. For this purpose, you will use the source file provided which already includes the data structures you will use, an implementation of the simple multiplication algorithm, as well as time measurement mechanisms. Your task is to fill in the `optMultiplication` function. While there are better ways to optimize this computation (e.g., using blocked matrix-matrix multiplication), you are expected to use the multi-pass matrix-vector multiplication idea discussed in class.
- b. [25 pts] *Performance Analysis:* Test the performance of your implementation for a set of different N (number of rows), M (number of columns) and B (blocking factor) values. Identify the different regimes as indicated by the relative speed up of the optimized implementation over the naïve one, and outline the conditions (in terms of matrix sizes, shapes, etc.) under which you observe these regimes.
Hint: Experiment with different matrix sizes (small to large) and input matrix shapes (short & wide rectangle, square, tall & skinny rectangle) while using a reasonable blocking factor.

- c. [25 pts] *Cache Performance Measurement*: Using a cache performance measurement tool (see instructions below on TAU) and your knowledge about how caches work, explain the root causes for the different regimes you observe in part b. For example, you can point to the ratio of cache misses you observe at various levels (L1, L2 or L3) using the naïve vs. optimized implementations to explain the relative performance differences.

Instructions:

- **Using the intel18 system:** You should use the intel18 system for all of your runs. Whenever you login to your hpcc you are in the gateway nodes and you should never run your programs on the gateway nodes. After logging in, run the following command to switch to the intel18 cluster (and then proceed with loading the necessary modules, if applicable)

```
ssh dev-intel18
```

- **Compiling your programs:** While you can use any compiler (different versions of GCC or Intel compiler), throughout this assignment please use GNU/6.4.0-2.28. GNU/6.4.0-2.28 together with OpenMPI 2.1.2 are necessary for compiling and running with TAU, a performance analysis tool that we will be using. So whenever you login to HPCC or in your (batch or interactive) queue runs, make sure to execute the following commands:

```
module unload GNU OpenMPI
module load GNU/6.4.0-2.28 OpenMPI/2.1.2
export TAU_SET_NODE=0
```

After loading the proper modules, you can compile the provided source file in one of the two ways below.

```
make multiplication
make multiplication-tau
```

Running the executables: After compilation, the provided source can be executed either in the "test" or in the "perf" modes on the intel18 dev-node:

```
./multiplication.x test inputfile
./multiplication.x perf N M B
```

where the *inputfile* contains a set of test runs. It starts with a line indicating the number of tests, followed by the tests themselves, each on a separate line. Each test is specified with N M B values.

N: number of matrix rows
M: number of matrix columns
B: blocking factor.

- We provide you a sample test file, which is the exact file that we will be using to test the accuracy (not performance) of your submission.
- "Test" mode is meant for accuracy testing only. You are expected to run your resulting implementation in the "perf" mode for timing and cache miss comparisons.
- **Measuring cache misses on HPCC:** For this purpose, we will be using the TAU performance analysis tool. Using hardware counters available through PAPI, TAU allows the measurement of an extensive list of CPU events. To compile your programs with TAU, follow the instructions below:

- i. Append the following lines to your Bash profile by editing the `.bashrc` file, which is a hidden file in your home directory (you can open it using vim or emacs, e.g. `vim ~/.bashrc`)

```
# PAPI
export PATH=/mnt/home/afibuzza/.local/papi/5.6.0/bin:$PATH
export INCLUDE_PATH=/mnt/home/afibuzza/.local/papi/5.6.0/include:$INCLUDE_PATH
export LD_LIBRARY_PATH=/mnt/home/afibuzza/.local/papi/5.6.0/lib:$LD_LIBRARY_PATH
export MANPATH=/mnt/home/afibuzza/.local/papi/5.6.0/man:$MANPATH

# TAU
export PATH=/mnt/home/afibuzza/.local/tau/2.28/x86_64/bin:$PATH
export INCLUDE_PATH=/mnt/home/afibuzza/.local/tau/2.28/include:$INCLUDE_PATH
export LD_LIBRARY_PATH=/mnt/home/afibuzza/.local/tau/2.28/x86_64/lib:$LD_LIBRARY_PATH
export TAU_MAKEFILE=/mnt/home/afibuzza/.local/tau/2.28/x86_64/lib/Makefile.tau-papi-mpi
export TAU_OPTIONS=-optCompInst
export TAU_METRICS=P_WALL_CLOCK_TIME:PAPI_L1_DCM:PAPI_L2_TCM:PAPI_L3_TCM
```

- ii. To see the complete list of available hardware counters, run on a dev-node

```
papi_avail
```

where supported counters have "Yes" in the third and fourth columns.

- iii. Now, compile the source file using the Tau MPI wrapper:

```
make multiplication-tau
```

- iv. You can now run the executable *instrumented* with TAU on a dev-node or using the queue system:

```
./multiplication-tau.x perf 1000 1000 200
```

- v. The performance data collected by TAU is saved automatically into default directories/files. Below are the directories/files you will see (as a result of the `TAU_METRICS` options specified):

```
MULTI__P_WALL_CLOCK_TIME/profile.0.0.0
```

```
MULTI__PAPI_L1_DCM/ profile.0.0.0
```

```
MULTI__PAPI_L2_TCM/profile.0.0.0
```

```
MULTI__PAPI_L3_TCM/profile.0.0.0
```

- vi. You can go to respective directories and then run the following command:

```
pprof > tauprofile
```

This will generate a `tauprofile` file in which you can see the cache misses in L1, L2 and L3 in a tabular view. Record these data.

- vii. You can simply view the `profile.0.0.0` files for each event using the `cat` command or a text editor (vim or emacs). The contents are simple enough to be understood in this way.
- viii. Note that, all the files generated by TAU are not overwritten, so make sure to record the cache miss count after each run and then remove/delete those folders before the next run to get the actual cache miss counts of a specific run.

- **Measuring your execution time and cache misses properly:** The wall-clock time measurement mechanism (based on the `gettimeofday()` function) implemented in the provided `main.c` file will allow you to measure the timings for a particular part of your program (see the skeleton code) precisely. However, on the dev-nodes there will be several other programs running simultaneously, and your measurements may not be very accurate due to the “background noise”. After making sure that your

program is bug-free and executes correctly (this can be done on the dev-intel18 node, for instance), a good way of getting reliable performance data for various input sizes is to use the interactive queue. **Please use the intel18 cluster for all your performance measurement runs!**

You can submit an interactive job request that will give you a dedicated node as follows:

```
salloc -N 1 -c 40 -C "intel18" --time=00:30:00 --mem=70G
```

or

```
salloc -N 1 -c 40 -C "skl" --time=00:30:00 --mem=70G
```

This will give you exclusive access to an intel18 node for 30 minutes. If you ask for a long job, your job may get delayed. **The default memory limit is set to be 750 MBs per job on HPCC systems, so it is very important that you ask for more as above.**

Once you are granted an interactive job, make sure to run your jobs one after the other (i.e., do not run them as background jobs, and do not worry about it if you do not know what background jobs mean). This is important because having multiple background jobs running simultaneously may create “noise” in the data you obtain.

Make sure to load the modules and make the binaries once you get an interactive job. Sometimes the interactive nodes will have a different module loaded which might not find the correct linker libraries.

- **Batch jobs:** Interactive jobs may sometimes be delayed significantly, and therefore insisting on this option may be counter-productive. For your convenience, we have provided a sample batch script named **job_script.sb**. Then submit the job using **sbatch job_script.sb**. Nice thing about batch jobs is that you can list multiple runs and automatically execute bunches of commands for copying/moving files around. See the scripts for examples.
- **Obtaining files from the git repo & submission:** You will clone the skeleton code and the testing input file through the instructor repository on the Gitlab server. Assuming that you have already cloned the instructor repository, you will need to pull the most recently committed files for HW2 and move them over into your individual repository.

Then complete the homework and submit it using your own personal repository. Your submission will include exactly two files (at the top level folder for HW2, not under any subdirectories such as src, etc.):

- o Your final optMultiplication.c file
- o A pdf file named “HW2_yourMSUNetID.pdf”, which contains your answers to questions in the assignment.